

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df=pd.read_csv("insurance (1).csv")
df
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [3]: df.shape
```

Out[3]: (1338, 7)

```
In [4]: df.describe()
```

```
Out[4]:
```


	age	bmi	children	charges
<b>count</b>	1338.000000	1338.000000	1338.000000	1338.000000
<b>mean</b>	39.207025	30.663397	1.094918	13270.422265
<b>std</b>	14.049960	6.098187	1.205493	12110.011237
<b>min</b>	18.000000	15.960000	0.000000	1121.873900
<b>25%</b>	27.000000	26.296250	0.000000	4740.287150
<b>50%</b>	39.000000	30.400000	1.000000	9382.033000
<b>75%</b>	51.000000	34.693750	2.000000	16639.912515
<b>max</b>	64.000000	53.130000	5.000000	63770.428010

```
In [5]: df.isnull().sum()
```

```
Out[5]: age      0  
sex      0  
bmi      0  
children  0  
smoker   0  
region   0  
charges  0  
dtype: int64
```


```
In [6]: sns.boxplot(df.bmi)
```

```
Out[6]: <AxesSubplot:xlabel='bmi'>
```




```
In [11]: sns.distplot(df.age)
print("skewness", df.age.skew())
```

skewness 0.05567251565299186



```
In [18]: sns.distplot(df.bmi)
df.bmi.mean()
```

Out[18]: 30.66339686098655




```
In [16]: df.bmi.median()
```

Out[16]: 30.4


```
In [19]: sns.violinplot(df.charges)
```

```
Out[19]: <AxesSubplot:xlabel='charges'>
```



```
In [20]: sns.scatterplot(x='bmi',y='charges',data=df)
```

```
Out[20]: <AxesSubplot:xlabel='bmi', ylabel='charges'>
```



## BMI for males and females

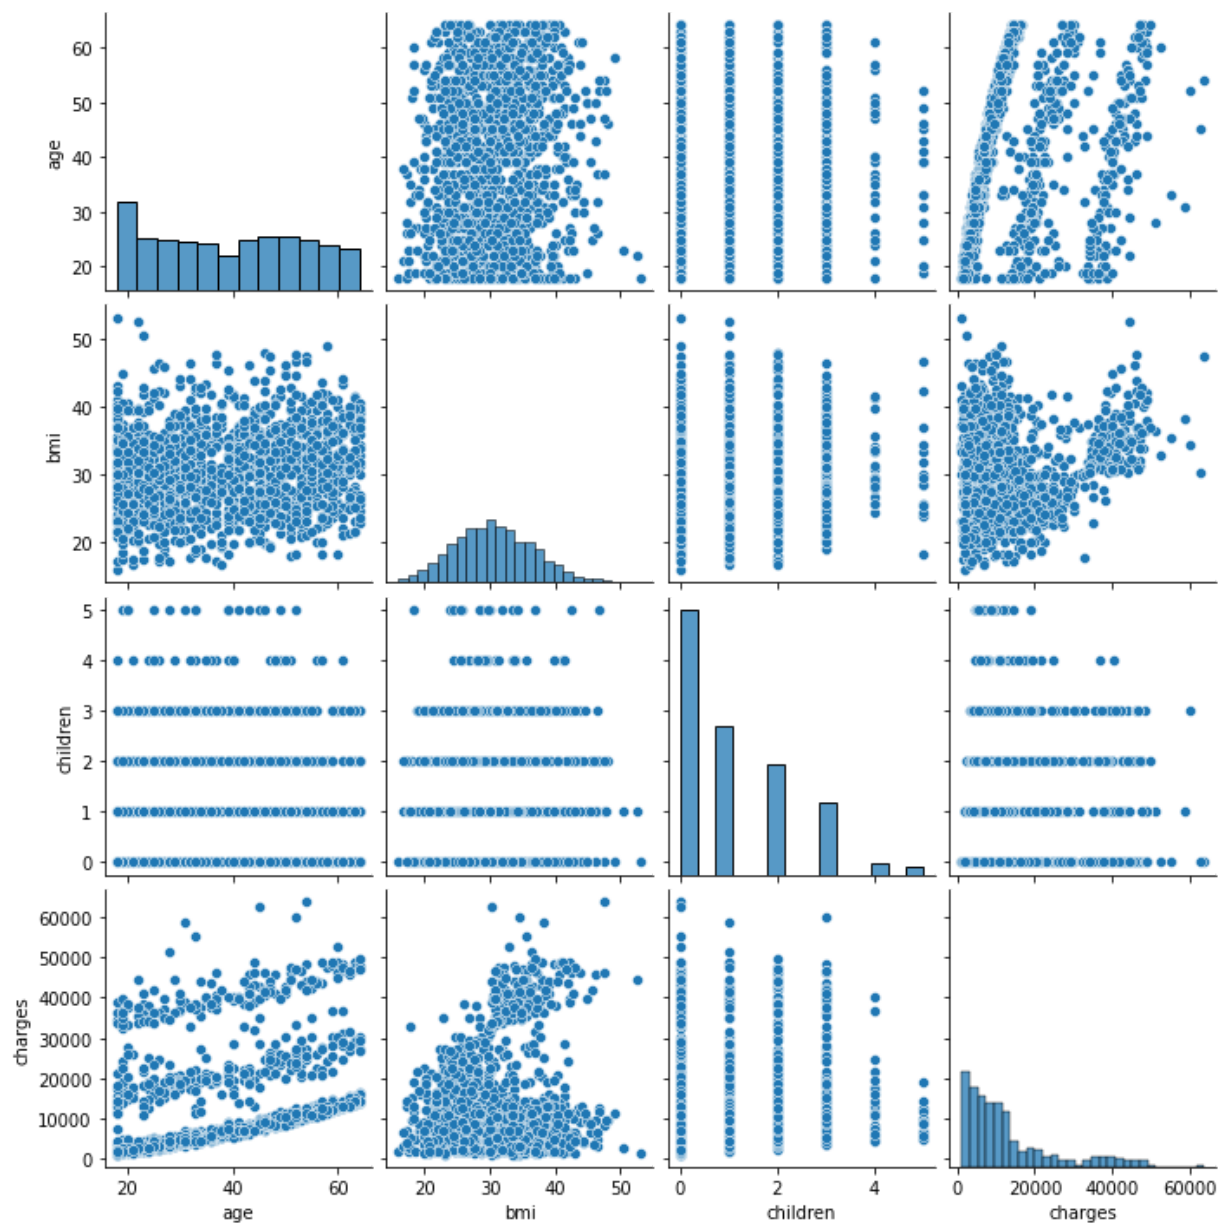
```
In [23]: df.groupby(['sex'])['bmi'].sum()
```

```
Out[23]: sex
female    20110.070
male      20917.555
Name: bmi, dtype: float64
```

```
In [26]: num_col=df.select_dtypes(include=['int','float'])
```

```
In [27]: sns.pairplot(data=num_col)
```


```
Out[27]: <seaborn.axisgrid.PairGrid at 0x1df04d51b70>
```



```
In [28]: corr=num_col.corr()
```

```
In [29]: sns.heatmap(corr,annot=True)
```

```
Out[29]: <AxesSubplot:>
```



## Handling Categorical data

```
In [30]: cat_col=df.select_dtypes(include=['object'])
```

In [31]: cat\_col

Out[31]:

	sex	smoker	region
0	female	yes	southwest
1	male	no	southeast
2	male	no	southeast
3	male	no	northwest
4	male	no	northwest
...	...	...	...
1333	male	no	northwest
1334	female	no	northeast
1335	female	no	southeast
1336	female	no	southwest
1337	female	yes	northwest

1338 rows × 3 columns

In [32]: `from sklearn.preprocessing import LabelEncoder`

In [33]: `le=LabelEncoder()`

In [34]: `for col in cat_col:  
 le=LabelEncoder()  
 df[col]=le.fit_transform(df[col])`

In [35]: `df.head()`

Out[35]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520

In [37]: `from sklearn.preprocessing import MinMaxScaler`

In [38]: `min_max=MinMaxScaler()`

```
In [39]: min_max.fit_transform(df[['bmi']])
```

```
Out[39]: array([[0.3212268 ],
 [0.47914985],
 [0.45843422],
 ...,
 [0.56201238],
 [0.26472962],
 [0.35270379]])
```

```
In [40]: from sklearn.preprocessing import StandardScaler
```

```
In [41]: ss=StandardScaler()
```

```
In [42]: ss.fit_transform(df[['bmi']])
```

```
Out[42]: array([[ -0.45332   ],
 [  0.5096211   ],
 [  0.38330685   ],
 ...,
 [  1.0148781   ],
 [-0.79781341   ],
 [-0.26138796   ]])
```

```
In [43]: df
```


```
Out[43]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520
...	...	...	...	...	...	...	...
1333	50	1	30.970	3	0	1	10600.54830
1334	18	0	31.920	0	0	0	2205.98080
1335	18	0	36.850	0	0	2	1629.83350
1336	21	0	25.800	0	0	3	2007.94500
1337	61	0	29.070	0	1	1	29141.36030

1338 rows × 7 columns



```
In [44]: for col in num_col:
print(col)
print('skewness',df[col].skew())
print('kurtosis',df[col].kurt())
plt.figure(figsize=(10,10))
sns.distplot(df[col])
plt.show()
```



```
bmi
skewness 0.2840471105987448
kurtosis -0.05073153135467834
```

```
In [45]: x=df.iloc[:, :-1]#Training data
y=df.iloc[:, -1]#Testing data
```

```
In [46]: x
```

Out[46]:

	age	sex	bmi	children	smoker	region
0	19	0	27.900	0	1	3
1	18	1	33.770	1	0	2
2	28	1	33.000	3	0	2
3	33	1	22.705	0	0	1
4	32	1	28.880	0	0	1
...	...	...	...	...	...	...
1333	50	1	30.970	3	0	1
1334	18	0	31.920	0	0	0
1335	18	0	36.850	0	0	2
1336	21	0	25.800	0	0	3
1337	61	0	29.070	0	1	1

1338 rows × 6 columns

```
In [47]: y
```

```
Out[47]: 0      16884.92400
          1      1725.55230
          2      4449.46200
          3      21984.47061
          4      3866.85520
          ...
          1333    10600.54830
          1334     2205.98080
          1335     1629.83350
          1336     2007.94500
          1337    29141.36030
          Name: charges, Length: 1338, dtype: float64
```

## splitting data

```
In [48]: from sklearn.model_selection import train_test_split
```

```
In [87]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=123)
```

```
In [88]: x_train.shape
```

```
Out[88]: (936, 6)
```

```
In [89]: y_train.shape
```

```
Out[89]: (936,)
```

```
In [90]: y_test.shape
```

```
Out[90]: (402,)
```

```
In [91]: x_test.shape
```

```
Out[91]: (402, 6)
```

## Model Building

```
In [92]: from sklearn.linear_model import LinearRegression
```

```
In [93]: reg=LinearRegression()
```

In [94]: `reg.fit(x_train,y_train)`

Out[94]: 

▼ LinearRegression

LinearRegression()

In [95]: `y_train_pred=reg.predict(x_train)`  
`y_test_pred=reg.predict(x_test)`

In [96]: `y_train_pred`

Out[96]: `array([ 8.14259855e+03, 1.44168619e+04, 7.19982210e+03, 3.28854161e+04,`  
 `2.94610576e+04, 1.04383624e+04, 3.17952432e+04, 7.86558258e+03,`  
 `3.88561982e+04, 1.02715278e+04, 4.25579616e+03, 7.35158232e+03,`  
 `7.02951770e+03, 1.00499786e+04, 5.57487564e+03, 3.78743167e+04,`  
 `3.70483135e+04, 1.32132015e+04, 1.49220421e+04, 2.61804889e+04,`  
 `1.04762009e+04, 1.11006090e+04, 1.06094913e+04, 8.64186275e+03,`  
 `3.75294316e+04, 1.76231376e+04, 1.63425435e+04, 3.91902128e+04,`  
 `5.02438054e+03, 5.79430522e+03, 5.12243288e+03, 1.21390583e+04,`  
 `2.29349955e+04, 1.77119121e+03, 6.85939923e+03, 3.01274749e+04,`  
 `1.20099602e+04, 1.04651895e+04, 9.75434388e+03, 2.49476358e+04,`  
 `9.91246384e+03, 1.38667638e+04, 8.63380105e+02, 1.70674866e+03,`  
 `1.37601982e+04, 1.07485625e+04, 1.15694035e+04, 1.06739405e+04,`  
 `6.08206166e+03, 1.02482210e+04, 8.16692564e+03, 9.21591756e+03,`  
 `6.91477038e+03, 7.87067780e+03, 5.85496101e+03, 9.99147365e+03,`  
 `9.20926898e+03, 1.13498902e+04, 3.88131556e+03, 3.45246685e+04,`  
 `7.29280792e+03, 4.18254175e+03, 9.17976421e+03, 1.50257276e+04,`  
 `3.93093866e+04, 1.65143394e+04, 2.98790702e+04, 1.50850951e+04,`  
 `7.46120896e+03, 3.47287170e+04, 3.72705966e+04, 3.35430651e+04,`  
 `1.16192690e+04, 8.95128026e+03, 6.38458070e+02, 1.28701172e+04,`  
 `2.42451640e+03, 7.82680047e+03, 2.54030421e+03, 2.21188470e+03,`

In [97]: `y_test_pred`

Out[97]: `array([ 1.55319667e+04, 9.51619844e+03, 2.83341303e+04, 5.36165725e+03,`  
 `1.14977191e+04, 1.12770325e+04, 2.84818292e+03, 2.22248819e+03,`  
 `4.45419112e+03, 8.81101145e+03, 8.71714119e+03, 1.35995111e+04,`  
 `1.33444996e+04, 2.99293763e+04, 1.37720291e+04, 3.10510278e+04,`  
 `1.50832570e+04, 1.72526684e+03, 3.36403616e+04, 3.67784399e+04,`  
 `3.34411743e+04, 4.11934013e+04, 3.17252957e+03, 1.01712572e+04,`  
`-5.35325945e+01, 9.59185567e+03, 1.64327524e+04, 6.57095072e+03,`  
 `1.47513939e+04, -5.19679604e+02, 9.28219974e+03, 1.64996937e+04,`  
 `5.29474627e+03, 3.76431146e+03, 6.93573633e+03, 1.14603044e+04,`  
 `1.77348674e+04, 7.02628562e+03, 9.83898680e+03, 1.16517139e+04,`  
 `7.12601669e+03, 7.95848427e+03, 5.41639869e+03, 3.18049796e+04,`  
 `7.37291174e+03, 2.38865923e+03, 1.31400244e+04, 1.39119932e+04,`  
 `1.46154279e+04, 9.43537877e+03, 1.28432468e+04, 1.11711223e+04,`  
 `2.73020882e+04, 1.15288421e+04, 1.42796256e+04, 6.31439797e+03,`  
 `3.19993033e+04, 2.82742869e+04, 3.85495351e+04, 6.55096839e+03,`  
 `1.64136464e+04, 9.57496936e+03, 3.84240192e+04, 2.84782849e+04,`  
 `1.75739994e+03, 4.26485403e+03, 1.78157402e+04, 5.71331001e+03,`  
 `1.55985407e+04, 1.48959207e+04, 1.08323710e+04, 1.16182336e+04,`  
 `1.63181998e+04, 1.68925411e+04, 1.15318732e+04, 3.19025444e+04,`  
 `4.12919789e+03, 1.45816675e+04, 5.67099084e+03, 3.11500552e+04,`  
 `2.01618635e+03, 3.11300110e+04, 2.80499493e+04, 1.53960828e+04,`  
 `2.55145446e+03, 1.64468305e+04, 2.91236269e+04, 1.32726759e+04,`  
 `1.04970262e+04, 3.91850034e+04, 9.76819203e+03, 9.30734124e+03,`  
 `8.73865648e+03, 6.27868550e+03, -7.55691365e+01, 2.70515287e+03,`  
 `6.32895650e+03, 8.20785574e+03, 4.14838122e+03, 3.90290499e+03,`  
 `1.13979293e+04, 8.84075362e+03, 1.33281031e+04, 6.52670837e+03,`  
 `7.53622608e+03, 3.43810181e+03, 2.76388992e+04, 8.35126293e+03,`  
 `1.35949601e+04, 8.61353611e+03, 4.54037386e+03, 3.73443835e+04,`  
 `3.16514992e+04, 1.38079406e+04, 3.11290203e+04, 1.35403313e+04,`  
 `7.95451984e+03, 9.47231094e+03, 1.30148719e+04, 7.36564336e+03,`  
 `3.83616842e+04, 2.67319761e+04, 5.35488794e+03, 1.15002556e+04,`  
 `5.60979051e+03, 4.12484655e+03, 3.30418557e+04, 7.58684918e+03,`  
 `3.74427474e+04, 3.67687976e+04, 4.50612920e+03, 6.30232908e+03,`  
`-2.87671664e+02, 1.45191650e+04, 9.38679336e+03, 6.90663700e+03,`  
 `6.22918881e+03, 5.86783454e+03, 4.06554105e+03, 1.37937875e+04,`  
 `4.73126922e+03, 2.38285347e+04, 2.78627153e+04, 8.72213751e+03,`  
 `2.58279899e+04, 9.59611791e+03, 3.03316405e+04, 7.17066802e+03,`  
 `8.43305395e+03, 2.09817518e+03, 3.27086839e+03, 1.26957328e+04,`  
 `2.61781984e+04, 1.17891682e+04, 3.47809859e+03, 5.49560114e+03,`  
 `3.75241902e+04, 1.11596733e+04, 1.35438013e+04, 5.87260471e+03,`  
 `5.81770770e+03, 1.05487716e+04, 8.93170315e+03, 3.63710003e+04,`  
 `3.04315825e+03, 3.85215143e+03, 1.18514795e+04, 2.86387039e+04,`  
 `2.70018427e+04, 1.07802734e+04, 5.70642357e+03, 3.00812469e+04,`  
 `1.21957550e+04, 1.45541785e+04, 1.17352598e+04, 5.64701752e+03,`  
 `1.84153532e+03, 9.94003315e+03, 1.29296072e+04, 7.39543611e+03,`  
 `7.65280796e+03, 1.17432700e+04, 9.94115748e+03, 4.10483896e+03,`  
 `1.19659554e+04, 7.30769482e+03, 1.40970913e+04, 4.13228080e+03,`  
 `8.03043840e+03, 2.88674077e+04, 3.17002102e+04, 1.35369887e+04,`  
 `3.39015307e+04, 1.06213223e+03, 1.03526887e+04, 8.09615765e+03,`  
 `2.32273879e+04, 9.25822986e+03, 2.97393237e+04, 1.27214707e+04,`  
 `3.90722165e+04, 1.36230439e+04, -2.99644103e+01, 7.47724058e+03,`  
 `4.28091561e+03, 2.73581031e+04, 3.48158698e+04, 3.89389329e+04,`  
 `3.09052244e+03, 2.93929193e+03, 9.98698836e+03, 1.15692876e+03,`  
 `1.26047380e+04, 1.00300914e+04, 1.24958198e+04, 3.86518251e+04,`

```

1.23996241e+04, 6.88255656e+03, 2.71786300e+04, 1.29194602e+04,
2.70165538e+04, 7.10766927e+03, 7.33134850e+03, 1.44750428e+04,
2.77779582e+04, 3.76285395e+03, 7.31137442e+03, 5.65445014e+02,
1.74769841e+04, 1.70913965e+04, 9.52736731e+03, 3.40518588e+03,
6.18511704e+03, 2.79176393e+04, 9.63166610e+03, 1.13832831e+04,
5.37366557e+03, 1.56342431e+04, 3.42469370e+04, 3.35410094e+04,
8.99690168e+03, 1.00333103e+04, 1.32924073e+04, 8.24691066e+03,
1.12495221e+04, 1.03973552e+04, 1.48425833e+04, 3.25539489e+03,
3.58806200e+03, 3.22130727e+04, 1.16420007e+04, 9.63386122e+03,
8.50238635e+03, 1.70087844e+04, 1.19765818e+04, 4.77634443e+03,
3.12379537e+04, 1.58230023e+04, 1.28070722e+04, 1.18455412e+04,
3.31769339e+04, 5.54078523e+03, 3.98884514e+04, 8.70147643e+03,
5.79285749e+03, 1.02641939e+04, 5.21058590e+03, 2.50289711e+04,
7.40295581e+03, 3.37059691e+04, 8.82486459e+03, 9.77548060e+03,
-1.25183968e+03, 1.14666295e+04, 9.18905376e+03, 9.16535474e+03,
2.68829516e+04, 1.23609424e+04, 3.24888257e+04, 9.78281450e+03,
1.12301545e+04, 3.76706271e+03, 3.17313066e+04, 9.96038934e+03,
8.63370438e+03, 1.18157339e+04, 3.29555778e+04, 1.23677871e+04,
1.53433236e+04, 5.66856492e+03, 3.19460949e+04, 3.53655239e+04,
1.08653539e+04, 9.61011205e+03, 1.61224852e+04, 3.68227046e+04,
2.55904946e+04, 1.00784976e+04, 1.59919112e+04, 7.55163584e+03,
2.60059837e+04, -1.92441073e+03, 1.13605277e+04, 3.03596335e+04,
1.12279593e+04, 1.00198491e+04, 1.69164038e+04, 3.40160568e+04,
3.29197564e+04, 6.29369987e+03, 2.74216779e+04, 2.61558021e+03,
8.30719989e+03, 2.93228025e+04, 1.25421749e+04, 2.78595452e+04,
8.57046993e+03, 2.31952511e+04, 9.96794609e+03, 1.04089920e+04,
1.42784508e+04, 7.89470163e+03, 9.21546664e+03, 1.32020577e+04,
5.48826723e+03, 1.26750736e+04, 1.64830856e+04, 1.00153872e+04,
1.55638486e+04, 1.57537691e+04, 1.31452956e+04, 1.45713142e+04,
1.38469288e+04, 1.01636814e+03, 7.18344205e+03, 2.85892454e+04,
5.36476201e+03, 3.22051578e+04, 4.16181223e+03, 1.86333759e+04,
3.63884713e+04, 1.18096382e+04, 5.33764195e+03, 6.22772750e+03,
1.48084929e+04, 6.18215813e+02, 9.91435021e+03, 1.19284672e+04,
4.57998857e+03, 1.27685313e+04, 7.73746093e+03, 2.98450936e+04,
1.61146659e+04, 7.08855680e+03, 6.78557149e+03, 1.05191398e+04,
3.86806571e+03, 1.91742185e+03, 4.66380723e+03, 2.83632844e+04,
3.21561348e+04, 7.03715628e+03, 7.67263787e+03, 1.31062757e+04,
3.03178154e+03, 6.64696133e+03, 7.35869825e+03, 1.10985059e+04,
9.39670439e+03, 1.09662730e+04, 7.98811625e+03, 4.50535011e+03,
1.01465310e+04, 3.38432888e+03, 2.66717036e+04, 3.70822566e+04,
2.90196249e+04, 7.00951038e+03, 2.72013736e+04, 6.40787940e+03,
9.55985936e+03, 2.60389015e+03, 4.48312769e+03, 1.28364091e+04,
1.57051610e+03, 1.81553709e+04, 6.50138072e+03, 1.32231656e+04,
4.10497994e+04, 1.28352915e+04, 2.07569283e+03, 1.67172249e+04,
9.30750559e+03, 1.61305249e+04, 1.07594164e+04, 3.50833178e+03,
3.58727811e+03, 7.03140088e+03, 5.69505670e+03, 2.80961794e+04,
7.25487388e+03, 4.92622786e+03])

```

In [98]: `reg.intercept_`

Out[98]: -11966.536839327007

```
In [99]: reg.coef_
```

```
Out[99]: array([ 258.01052077, -89.42500591,  343.81965818,  479.31668732,  
                23703.74833645, -388.43417176])
```

## Find R2 Score

```
In [100]: from sklearn.metrics import r2_score, mean_squared_error
```

```
In [101]: train_r2_score = r2_score(y_train, y_train_pred)
```

```
In [102]: test_r2_score = r2_score(y_test, y_test_pred)
```

```
In [103]: train
```

```
Out[103]: 0.7454618032357492
```

```
In [104]: test
```

```
Out[104]: 0.7626072475418817
```

```
In [105]: def metric(y_actual, y_pred):  
            r2 = r2_score(y_actual, y_pred)  
            RMSE = np.sqrt(mean_squared_error(y_actual, y_pred))  
            print("r2 score: {} | RMSE: {}".format(round(r2, 2), round(RMSE, 2)))
```

```
In [106]: print("Training performance")  
            metric(y_train, y_train_pred)  
            print("Testing performance")  
            metric(y_test, y_test_pred)
```

```
Training performance  
r2 score: 0.75 | RMSE: 6140.17  
Testing performance  
r2 score: 0.76 | RMSE: 5823.57
```

```
In [ ]:
```

```
In [ ]:
```